

## Assignment 7 (Sample problems and solutions)

### Object-Relational Database Programming

This is an assignment to test your understanding of Lecture 19 on Object-Relational Database programming.

You should submit a .sql file with the solutions and a separate .txt file with appropriate outputs.

#### 1. For this problem you can not use arrays.

Consider the relational schema `Tree(parent int, child int)` representing the schema for storing a rooted tree  $T$ .<sup>1</sup> A pair of nodes  $(m, n)$  is in `Tree` if  $m$  is the parent of  $n$  in  $T$ . Notice that a node  $m$  can be the parent of multiple children but a node  $n$  can have at most one parent node.

It should be clear that for each pair of different nodes  $m$  and  $n$  in  $T$ , there is a unique shortest path of nodes  $(n_1, \dots, n_k)$  in  $T$  from  $m$  to  $n$  provided we interpret the edges in  $T$  as undirected. A good way to think about this path from a node  $m$  to a node  $n$  is to first consider the *lowest common ancestor node* of  $m$  of  $n$  in  $T$ . Then the unique path from  $m$  to  $n$  is the path that is comprised of the path up the tree from  $m$  to this common ancestor and then, from this common ancestor, the path down the tree to the node  $n$ . (Note that in this path  $n_1 = m$  and  $n_k = n$ .)

Define the *distance* from  $m$  to  $n$  to be  $k - 1$  if  $(n_1, \dots, n_k)$  is the unique shortest path from  $m$  to  $n$  in  $T$ .

Write a PostgreSQL function `distance(m,n)` that computes the distance in  $T$  for any possible pair  $m$  and  $n$  in  $T$ .<sup>2</sup>

For example, if  $m$  is the parent of  $n$  in  $T$  then `distance(m,n)` = 1 because the shortest path from  $m$  to  $n$  is  $(m,n)$  which has length 1. If  $m$  is the grandparent of  $n$  in  $T$  then `distance(m,n)` = 2 since  $(m,p,n)$  is the path from  $m$  to  $n$  where  $p$  is the parent of  $m$  and  $p$  is a child of  $n$ . And if  $m$  and  $n$  have a common grandparent  $k$  then `distance(m,n)` = `distance(m,k)` + `distance(k,n)` = 4, etc.

---

<sup>1</sup>We assume that a tree is a connected graph with a finite number of nodes.

<sup>2</sup>Incidentally, if  $m = n$  then `distance(m,n)` = 0 since, in this case, the unique path from  $m$  to  $n$  is  $(m)$  which is a path of length 0.

2. In this problem, you can not use arrays.

Consider the following relational schemas. (You can assume that the domain of each of the attributes in these relations is `int`.)

`partSubpart(pid,sid,quantity)`  
`basicPart(pid,weight)`

A tuple  $(p, s, q)$  is in `partSubPart` if part  $s$  occurs  $q$  times as a **direct** subpart of part  $p$ . For example, think of a car  $c$  that has 4 wheels  $w$  and 1 radio  $r$ . Then  $(c, w, 4)$  and  $(c, r, 1)$  would be in `partSubpart`. Next think of a wheel  $w$  that has 5 bolts  $b$ . Then  $(w, b, 5)$  would be in `partSubpart`.

A tuple  $(p, w)$  is in `basicPart` if basic part  $p$  has weight  $w$ . A basic part is defined as a part that does not have subparts. In other words, the pid of a basic part does not occur in the pid column of `partSubpart`.

(In the above example, a bolt and a radio would be basic parts, but car and wheel would not be basic parts.)

We define the *aggregated weight* of a part inductively as follows:

- (a) If  $p$  is a basic part then its aggregated weight is its weight as given in the `basicPart` relation
- (b) If  $p$  is not a basic part, then its aggregated weight is the sum of the aggregated weights of its subparts, each multiplied by the quantity with which these subparts occur in the `partSubpart` relation:

**Example:** The following example is based on a desk lamp with pid 1. Suppose a desk lamp consists of 4 bulbs (with pid 2) and a frame (with pid 3), and a frame consists of a post (with pid 4) and 2 switches (with pid 5). Furthermore, we will assume that the weight of a bulb is 5, that of a post is 50, and that of a switch is 3.

Then the `partSubpart` and `basicPart` relation would be as follows:

partSubPart			basicPart	
pid	sid	quantity	pid	weight
1	2	4	2	5
1	3	1	4	50
3	4	1	5	3
3	5	2		

Then the aggregated weight of a lamp is  $4 \times 5 + 1 \times (1 \times 50 + 2 \times 3) = 76$ .

Write a PostgreSQL function `aggregatedWeight(p integer)` that returns the aggregated weight of a part  $p$ .

- (a) A problem very similar to the aggregated weight problem is described in the PostgreSQL manual page  
`https://www.postgresql.org/docs/8.4/queries-with.html`  
Adapt the recursive program in that manual page to the aggregated weight problem.
- (b) Write a non-recursive version for aggregate weight problem. Obviously, your program will require looping statements.

3. **In this problem, you can use arrays, but only as a mechanism to represents subsets of  $A(x)$ .**

Consider the relation schema  $A(x)$  representing a schema for storing a set  $A$ . (You can assume that the domain of attribute  $x$  is integer.)

Using arrays to represent sets, write a PostgreSQL program

`superSetsOfSet(X int[])`

that returns a relation that contains each subset of  $A$  that is a superset of  $X$ , i.e., each set  $Y$  such that  $X \subseteq Y \subseteq A$ .

For example, if  $X = \{\}$ , then `superSetsofSets(X)` should return the relation that consist of each element in the powerset of  $A$ .

4. **In this problem, you can use arrays, but only as a mechanism to represents sets of words.**

Consider the relation schema `document(doc int, words text[])` representing a relation of pairs  $(d, W)$  where  $d$  is a unique document id and  $W$  denotes the set of words that occur in  $d$ .

Let  $\mathbf{W}$  denote the set of all words that occur in the documents and let  $t$  be a positive integer denoting a *threshold*.

Let  $X \subseteq \mathbf{W}$ . We say that  $X$  is  $t$ -frequent if

$$\text{count}(\{d | (d, W) \in \text{document and } X \subseteq W\}) \geq t$$

In other words,  $X$  is  $t$ -frequent if there are at least  $t$  documents that contain all the words in  $X$ .

Write a PostgreSQL program `frequentSets(t int)` that returns the relation of all  $t$ -frequent sets.

In a good solution for this problem, you should use the following rule: if  $X$  is not  $t$ -frequent then any set  $Y$  such that  $X \subseteq Y$  is not  $t$ -frequent either. In the literature, this is called the *Apriori* rule of the frequent itemset mining problem. This rule allows you to avoid examining supersets of sets that are not frequent. This can drastically reduce the search space.